

George Suarez
David Cruz
CSE 460

Lab 9 – Page Replacement Algorithms and XV6 Priority Scheduling

1) First-in First-out (FIFO) Replacement

- Compile and Execute *fifo1.cpp*

The following is a sample input and output of this program:

```
-----  
Enter max. number of frames allowed in main memory: 3  
  
Enter sequence of page requests (-99 to terminate).  
New page : 2  
  
page 2 is allocated to frame 0  
Total page faults = 1  
New page : 3  
  
page 3 is allocated to frame 1  
Total page faults = 2  
New page : 2  
  
page 2 already in frame 0  
New page : -99  
  
Total number of faults: 2  
-----
```

- Try the Belady's anomaly examples discussed in class. Did you observe the Belady's anomaly?

Yes. When all the frames are used up, the oldest page has been replaced with a newer page. Also, the number of faults depends on what the maximum of frames has been created, and if a page is already in a certain frame then the page fault is not incremented at all.

georgesuarez at MacBook-Pro in ~/University/CSE-460/Labs/Lab 9 on master*

\$./fifo1

Enter max. number of frames allowed in main memory: 4

Enter sequence of page requests (-99 to terminate).

New page : 0

page 0 is allocated to frame 0

Total page faults = 1

New page : 1

page 1 is allocated to frame 1

Total page faults = 2

New page : 2

page 2 is allocated to frame 2

Total page faults = 3

New page : 3

page 3 is allocated to frame 3

Total page faults = 4

New page : 0

page 0 already in frame 0

New page : 1

page 1 already in frame 1

New page : 4

page 4 is allocated to frame 0

Total page faults = 5

New page : 0

page 0 is allocated to frame 1

Total page faults = 6

New page : 1

page 1 is allocated to frame 2

Total page faults = 7

New page : 2

page 2 is allocated to frame 3

Total page faults = 8

New page : 3

page 3 is allocated to frame 0

Total page faults = 9

New page : 4

page 4 is allocated to frame 1

Total page faults = 10

New page : -99

Total number of faults: 10

2) Multithreads for FIFO Program

- Implement displayMsg.cpp. Run displayMsg in one X-term and then fifo2 in another X-term. Repeat the examples of Belady's anomaly discussed above.

Fifo2.cpp

```
#include <SDL2/SDL.h>
```

```
#include <SDL2/SDL_thread.h>
```

```
#include <sys/ipc.h>
```

```
#include <sys/msg.h>
```

```
#include <sys/types.h>
```

```
#include <stdio.h>
```

```
#include <stdlib.h>
```

```

#include <math.h>
#include <signal.h>
#include <unistd.h>
#include <iostream>
#include <deque>

using namespace std;

class Cframe
{
public:
    int frameNo;    // frame number
    int pageNo;     // page number
    int r;          // reference bit
    Cframe(int n, int p) // constructor
    {
        frameNo = n;
        pageNo = p; // no page loaded at beginning
        r = 0;
    }
};

deque <Cframe> Q;
int nFaults = 0;
int page, frame;
SDL_mutex *mutex1;
SDL_cond *updateQueue; //condition variable
bool update = false;
bool quit = false;

```

```
#define MAX_TEXT 512
```

```
struct my_msg_st
```

```
{
```

```
    long int my_msg_type;
```

```
    char some_text[MAX_TEXT];
```

```
};
```

```
int displayMsg(void *data)
```

```
{
```

```
    struct my_msg_st some_data;
```

```
    int msgid;
```

```
    char buffer[BUFSIZ];
```

```
msgid = msgget((key_t)1234, 0666 | IPC_CREAT);
```

```
if (msgid == -1)
```

```
{
```

```
    fprintf(stderr, "msgget failed with error: %d\n", errno);
```

```
    exit(EXIT_FAILURE);
```

```
}
```

```
while (true)
```

```
{
```

```
    SDL_LockMutex(mutex1);
```

```
    while (!update && !quit)
```

```
        SDL_CondWait(updateQueue, mutex1);
```

```
    update = false;
```

```
    SDL_LockMutex(mutex1);
```

```
    sprintf(buffer, "%d,%d,%d\n", page, frame, nFaults);
```

```
    some_data.my_msg_type = 1;
```

```

strcpy(some_data.some_text, buffer);

if (msgsnd(msgid, (void *)&some_data, MAX_TEXT, 0) == -1)
{
    fprintf(stderr, "msgsnd failed\n");
    exit(EXIT_FAILURE);
}
if (page == -99)
    break;
}
exit(EXIT_SUCCESS);
}

```

```

void fault()
{
    nFaults++;
}

int search(const deque<Cframe> &q, int p)
{
    int n = q.size();
    for (int i = 0; i < n; i++)
    {
        if (q[i].pageNo == p)
            return q[i].frameNo;
    }
    return -1;
}

```

```

int main()
{

```

```
SDL_Thread *tid = SDL_CreateThread(displayMsg, "Send Thread", (char *) "Send-thread");
```

```
int maxFrames;
```

```
cout << "\nEnter max. number of frames allowed in main memory: ";
```

```
cin >> maxFrames;
```

```
int n;
```

```
cout << "Enter sequence of page requests (-99 to terminate).\n";
```

```
while (true)
```

```
{
```

```
    cout << "New page : ";
```

```
    cin >> page;
```

```
    if (page == -99)
```

```
    {
```

```
        quit = true;
```

```
        SDL_CondSignal(updateQueue);
```

```
        break;
```

```
    }
```

```
    if ((frame = search(Q, page)) != -1)
```

```
    {
```

```
        ;
```

```
    }
```

```
    else
```

```
    {
```

```
        n = Q.size();
```

```
        if (n < maxFrames)
```

```
        {
```

```
            Cframe aFrame(n, page);
```

```
            Q.push_back(aFrame);
```

```

        frame = aFrame.frameNo;
    }
    else
    {
        Cframe aFrame = Q.front();
        Q.pop_front();
        aFrame.pageNo = page;
        Q.push_back(aFrame);
        frame = aFrame.frameNo;
    }
    fault();
}

SDL_LockMutex(mutex1);
update = true;
SDL_CondSignal(updateQueue);
SDL_UnlockMutex(mutex1);
}

SDL_WaitThread(tid, NULL);
return 0;
}

```

displayMsg.cpp

```

include <errno.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <sys/ipc.h>
#include <sys/msg.h>
#include <sys/types.h>

```



```

#include <unistd.h>
#include <iostream>

#define MAX_TEXT 512

struct my_msg_st
{
    long int my_msg_type;
    char some_text[MAX_TEXT];
};

int main()
{
    int running = 1;
    struct my_msg_st some_data;
    int msgid, page, frame, faults;
    msgid = msgget((key_t)1234, 0666 | IPC_CREAT);
    if (msgid == -1)
    {
        fprintf(stderr, "msgget failed with error: %d\n", errno);
        exit(EXIT_FAILURE);
    }
    printf("Page\tFrame\tTotal Faults\n");
    while (1)
    {
        if (msgrcv(msgid, (void *)&some_data, MAX_TEXT, 0, 0) == -1)
        {
            fprintf(stderr, "msgrcv failed with error: %d\n", errno);
            exit(EXIT_FAILURE);
        }
    }
}

```

```

sscanf(some_data.some_text, "%d,%d,%d\n", &page, &frame, &faults);
if (strcmp(some_data.some_text, "-99", 3) == 0)
{
    printf("\nTerminal ending...\n");
    running = 0;
    break;
}
printf("%d\t%d\t%d\n", page, frame, faults);
}
if (msgctl(msgid, IPC_RMID, 0) == -1)
{
    fprintf(stderr, "msgctl failed with error: %d\n", errno);
    exit(EXIT_FAILURE);
}
exit(EXIT_SUCCESS);
}

```

georgesuarez at MacBook-Pro in ~/University/CSE-460/Labs/Lab 9 on master*

\$./fifo2

Enter max. number of frames allowed in main memory: 3

Enter sequence of page requests (-99 to terminate).

New page : 0

New page : 1

New page : 2

New page : 3

New page : 0

New page : 1

New page : 4

New page : 0

New page : 4
New page : 3
New page : 2
New page : 1
New page : 2
New page : 3
New page : 4
New page : -99

georgesuarez at MacBook-Pro in ~/University/CSE-460/Labs/Lab 9 on master*

\$./displayMsg

Page	Frame	Total Faults
0	0	1
1	1	2
2	2	3
3	0	4
0	1	5
1	2	6
4	0	7
0	1	7
4	0	7
3	1	8
2	2	9
1	0	10
2	2	10
3	1	10
4	1	11

Terminal ending...

3) Implement One of the following, Second Chance or LRU:

a. Second Chance

Modify fifo2.cpp to fifo3.cpp to implement the second-chance FIFO replacement discussed above. Compare the total faults for this algorithm and those of fifo2.cpp. Which one yields better results?

fifo2 produces slightly better results since a total of 11 faults were created as opposed to the 12 faults created in fifo3.

Fifo3.cpp

```
#include <iostream>
#include <SDL2/SDL.h>
#include <SDL2/SDL_thread.h>
#include <stdio.h>
#include <stdlib.h>
#include <sys/types.h>
#include <sys/msg.h>
#include <unistd.h>
#include <deque>
using namespace std;
class Cframe
{
public:
    int frameNo;
    // frame number
    int pageNo;
    // page number
    int r;
    // reference bit
    Cframe(int n, int p)
    {
```

```

        frameNo = n;
        pageNo = p;
        // no page loaded at beginning
        r = 0;
    }
};

deque<Cframe> Q;

int nFaults = 0;
int page, frame;

SDL_mutex *mutex1;
SDL_cond *updateQueue; // condition variable

bool update = false;
bool quit = false;

#define MAX_TEXT 512

struct my_msg_st
{
    long int my_msg_type;
    char some_text[MAX_TEXT];
};

int displayMsg(void *data)
{
    struct my_msg_st some_data;
    int msgid;
    char buffer[BUFSIZ];
    msgid = msgget((key_t)1234, 0666 | IPC_CREAT);
    if (msgid == -1)
    {
        fprintf(stderr, "msgget failed with error: %d\n", errno);
        exit(EXIT_FAILURE);
    }

```

```

}
while (true)
{
    SDL_LockMutex(mutex1);
    while (!update && !quit)
    {
        SDL_CondWait(updateQueue, mutex1);
    }
    update = false;
    SDL_LockMutex(mutex1);
    sprintf(buffer, "%d,%d,%d\n", page, frame, nFaults);
    some_data.my_msg_type = 1;
    strcpy(some_data.some_text, buffer);
    if (msgsnd(msgid, (void *)&some_data, MAX_TEXT, 0) == -1)
    {
        fprintf(stderr, "msgsnd failed\n");
        exit(EXIT_FAILURE);
    }
    if (page == -99)
    {
        break;
    }
}
exit(EXIT_SUCCESS);
}

void fault()
{
    nFaults++;
}

int search(deque<Cframe> &q, int p)

```

```

{
    int n = q.size();
    for (int i = 0; i < n; i++)
    {
        if (q[i].pageNo == p)
        {
            q[i].r = 1;
            // recently referenced
            return q[i].frameNo;
        }
    }
    return -1;
}

```

```

int main()
{
    SDL_Thread *tid = SDL_CreateThread(displayMsg, "Send Message" ,(char *)"Send-
thread");

    int maxFrames;
    cout << "\nEnter max. number of frames allowed in main memory: ";
    cin >> maxFrames;

    int n;
    cout << "Enter sequence of page requests (-99 to terminate).\n";
    while (true)
    {
        cout << "New page: ";
        cin >> page;
        if (page == -99)
        {
            quit = true;
            SDL_CondSignal(updateQueue);

```

```

        break;
    }
    if ((frame = search(Q, page)) != -1)
    {
        ;
    }
    else
    {
        n = Q.size();
        if (n < maxFrames)
        {
            Cframe aFrame(n, page);
            Q.push_back(aFrame);
            frame = aFrame.frameNo;
        }
        else
        {
            Cframe aFrame = Q.front();
            while (aFrame.r == 1)
            {
                // find oldest page that
                Q.pop_front();
                // has r == 0; set all r
                aFrame.r = 0;
                // flags to 0 until one is
                Q.push_back(aFrame);
                // found
                aFrame = Q.front();
            }
            Q.pop_front();
            aFrame.pageNo = page;

```



```

        Q.push_back(aFrame);
        frame = aFrame.frameNo;
    }
    fault();
}
SDL_LockMutex(mutex1);
update = true;
SDL_CondSignal(updateQueue);
SDL_UnlockMutex(mutex1);
}
SDL_WaitThread(tid, NULL);
return 0;
}

```

georgesuarez at MacBook-Pro in ~/University/CSE-460/Labs/Lab 9 on master*

\$./fifo3

Enter max. number of frames allowed in main memory: 3

Enter sequence of page requests (-99 to terminate).

New page: 0

New page: 1

New page: 2

New page: 3

New page: 0

New page: 1

New page: 4

New page: 0

New page: 4

New page: 3

New page: 2

New page: 1

New page: 2

New page: 3

New page: 4

New page: -99

\$./displayMsg

Page	Frame	Total Faults
0	0	1
1	1	2
2	2	3
3	0	4
0	1	5
1	2	6
4	0	7
0	1	7
4	0	7
3	2	8
2	1	9
1	2	10
2	1	10
3	0	11
4	2	12

Terminal ending...

4) XV6 Process Priority

1. Giving high priority to a newly loaded process by adding a *priority* statement in *exec.c*:

```
...  
curproc->tf->eip = elf.entry; // main  
curproc->tf->esp = sp;  
curproc->priority = 2;      // Added statement  
switchvm(curproc);  
freevm(oldpgdir);
```

2. Modifying *foo.c* so that the parent waits for the children:

```
for ( k = 0; k < n; k++ ) {  
    id = fork();  
    if ( id < 0 ) {  
        printf(1, "%d failed in fork!\n", getpid() );  
    } else if ( id > 0 ) { // parent  
        //printf(1, "Parent %d creating child %d\n", getpid(), id );  
        wait();  
    } else { // child  
        //printf(1, "Child %d created\n", getpid() );  
        for ( z = 0; z < 8000000.0; z += 0.001 )  
            x = x + 3.14 * 89.64; // useless calculations to consume CPU time  
        break;  
    }  
}
```

3. Observing the default round-robin (RR) scheduling.

```
$ foo &; foo &; foo &
```

```
$ ps
```

	name	pid	state	priority
init	1		SLEEPING	2

sh	2	SLEEPING	2
foo	9	RUNNING	10
foo	8	SLEEPING	2
foo	5	SLEEPING	2
foo	7	SLEEPING	2
foo	10	RUNNABLE	10
foo	11	RUNNABLE	10
ps	12	RUNNING	2

\$ ps

name	pid	state	priority
init	1	SLEEPING	2
sh	2	SLEEPING	2
foo	9	RUNNABLE	10
foo	8	SLEEPING	2
foo	5	SLEEPING	2
foo	7	SLEEPING	2
foo	10	RUNNING	10
foo	11	RUNNABLE	10
ps	13	RUNNING	2

\$ ps

name	pid	state	priority
init	1	SLEEPING	2
sh	2	SLEEPING	2
foo	9	RUNNING	10
foo	8	SLEEPING	2
foo	5	SLEEPING	2
foo	7	SLEEPING	2
foo	10	RUNNABLE	10
foo	11	RUNNABLE	10

ps 14 RUNNING 2

4. Implementing Priority Scheduling in *proc.c*:

```
#define NULL 0

void
scheduler(void)
{
    struct proc *p;
    struct proc *p1;
    struct cpu *cpu = mycpu();
    cpu->proc = 0;

    for(;;){
        // Enable interrupts on this processor.
        sti();

        struct proc *highP = NULL;
        // Loop over process table looking for process to run.
        acquire(&ptable.lock);
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->state != RUNNABLE)
                continue;
            highP = p;
            // choose one with highest priority
            for(p1 = ptable.proc; p1 < &ptable.proc[NPROC]; p1++){
                if(p1->state != RUNNABLE)
                    continue;
                if ( highP->priority > p1->priority ) // larger value, lower priority
                    highP = p1;
            }
        }
    }
}
```

```

// Switch to chosen process. It is the process's job
// to release ptable.lock and then reacquire it
// before jumping back to us.
p = highP;
cpu->proc = p;
switchvm(p);
p->state = RUNNING;
// cprintf("Process %s with pid %d running\n with createTime %d\n", p->name, p->pid, p-
>createTime);

swtch(&(cpu->scheduler), p->context);
switchkvm();

// Process is done running for now.
// It should have changed its p->state before coming back.
cpu-> proc = 0;
}
release(&ptable.lock);

}
}

```

5. Observing the priority scheduling. We run xv6 with the scheduler and again use *foo* and *ps* to see how it works. We use *nice* to change the priority of a process.

```
$ foo &; foo &; foo &
```

```
$ ps
```

name	pid	state	priority
init	1	SLEEPING	2
sh	2	SLEEPING	2
foo	9	RUNNABLE	10

foo	8	SLEEPING	2
foo	5	SLEEPING	2
foo	7	SLEEPING	2
foo	10	RUNNING	10
foo	11	RUNNABLE	10
ps	12	RUNNING	2

\$ nice 11 8

\$ ps

name	pid	state	priority
init	1	SLEEPING	2
sh	2	SLEEPING	2
foo	9	RUNNABLE	10
foo	8	SLEEPING	2
foo	5	SLEEPING	2
foo	7	SLEEPING	2
foo	10	RUNNABLE	10
foo	11	RUNNING	8
ps	14	RUNNING	2

\$ ps

name	pid	state	priority
init	1	SLEEPING	2
sh	2	SLEEPING	2
foo	9	RUNNABLE	10
foo	8	SLEEPING	2
foo	5	SLEEPING	2
foo	7	SLEEPING	2
foo	10	RUNNABLE	10
foo	11	RUNNING	8

ps 15 RUNNING 2

Discussion: We did finished everything in this lab so we would give ourselves 20/20 points.