

George Suarez  
David Cruz  
CSE 460

## Lab 10 - Exploring XV6 File System

1. Examining the header file *fs.h* and the program *fs.c* that implements the file system.

*fs.h*

```
/* On-disk file system format.
// Both the kernel and user programs use this header file.

#define ROOTINO 1 // root i-number
#define BSIZE 512 // block size

// Disk layout:
// [ boot block | super block | log | inode blocks |
//                                free bit map | data blocks]
//
// mkfs computes the super block and builds an initial file system. The
// super block describes the disk layout:
struct superblock {
    uint size;           // Size of file system image (blocks)
    uint nblocks;        // Number of data blocks
    uint ninodes;        // Number of inodes.
    uint nlog;           // Number of log blocks
    uint logstart;       // Block number of first log block
    uint inodestart;     // Block number of first inode block
    uint bmapstart;      // Block number of first free map block
};

#define NDIRECT 12
#define NINDIRECT (BSIZE / sizeof(uint))
#define MAXFILE (NDIRECT + NINDIRECT)

// On-disk inode structure
struct dinode {
    short type;          // File type
    short major;         // Major device number (T_DEV only)
    short minor;         // Minor device number (T_DEV only)
    short nlink;         // Number of links to inode in file system
    uint size;           // Size of file (bytes)
    uint addrs[NDIRECT+1]; // Data block addresses
};

// Inodes per block.
#define IPB (BSIZE / sizeof(struct dinode))

// Block containing inode i
#define IBLOCK(i, sb) ((i) / IPB + sb.inodestart)

// Bitmap bits per block
#define BPB (BSIZE*8)

// Block of free map containing bit for block b
#define BBLOCK(b, sb) (b/BPB + sb.bmapstart)

// Directory is a file containing a sequence of dirent structures.
#define DIRSIZ 14

struct dirent {
    ushort inum;
    char name[DIRSIZ];
};
```

## *fs.c*

```
/* File system implementation. Five layers:
// + Blocks: allocator for raw disk blocks.
// + Log: crash recovery for multi-step updates.
// + Files: inode allocator, reading, writing, metadata.
// + Directories: inode with special contents (list of other inodes!)
// + Names: paths like /usr/rtn/xv6/fs.c for convenient naming.
//
// This file contains the low-level file system manipulation
// routines. The (higher-level) system call implementations
// are in sysfile.c.

#include "types.h"
#include "defs.h"
#include "param.h"
#include "stat.h"
#include "mmu.h"
#include "proc.h"
#include "spinlock.h"
#include "sleeplock.h"
#include "fs.h"
#include "buf.h"
#include "file.h"

#define min(a, b) ((a) < (b) ? (a) : (b))
static void itrunc(struct inode*);
// there should be one superblock per disk device, but we run with
// only one device
struct superblock sb;

// Read the super block.
void
readsb(int dev, struct superblock *sb)
{
    struct buf *bp;

    bp = bread(dev, 1);
    memmove(sb, bp->data, sizeof(*sb));
    brelse(bp);
}

// Zero a block.
static void
bzero(int dev, int bno)
{
    struct buf *bp;

    bp = bread(dev, bno);
    memset(bp->data, 0, BSIZE);
    log_write(bp);
    brelse(bp);
}

...

```

2. Add to a system call named `pfs()` in `fs.c` that prints out information of the file system

a. Adding the name to *syscall.h*

```
#define SYS_mkdir    20
#define SYS_close    21
#define SYS_cps      22
#define SYS_chpr     23
#define SYS_pfs      24
```

b. Adding the function prototype to *defs.h*

```
int      namecmp(const char*, const char*);
struct inode* namei(char*);
struct inode* nameiparent(char*, char*);
int      readi(struct inode*, char*, uint, uint);
void     stati(struct inode*, struct stat*);
int      writei(struct inode*, char*, uint, uint);
int      pfs(void);
```

- c. Added the function prototype to *users.h*

```
int sleep(int);
int uptime(void);
int cps ( void );
int chpr (int pid, int priority);
int pfs(void);
```

- d. Added the function call to *syscall.c*

```
int
sys_cps ( void )
{
    return cps ();
}

int
sys_chpr ( void )
{
    int pid, pr;
    if (argint(0, &pid) < 0)
        return -1;
    if (argint(1, &pr) < 0)
        return -1;

    return chpr ( pid, pr );
}

int sys_pfs(void)
{
    return pfs();
}
```

- e. Added the call to *usys.S*

```
SYSCALL(sbrk)
SYSCALL(sleep)
SYSCALL(uptime)
SYSCALL(cps)
SYSCALL(chpr)
SYSCALL(pfs)
--
```

- f. Added the call to *syscall.c*

```
extern int sys_chpr(void);
extern int sys_pfs(void);
[SYS_cps]    sys_cps,
[SYS_chpr]   sys_chpr,
[SYS_pfs]    sys_pfs,
```

- g. Add *pfs()* to *fs.c*

```
int
pfs()
{
    cprintf("sb: size  nblocks  ninodes nlog logstart inodestart bmap-start Inodes-per-block Bitmap-bits-per-block\n");
    cprintf("    %d\t    %d\t    %d    %d    %d \t %d \t %d \t %d \t %d \t %d\n", sb.size, sb.nblocks,
            sb.ninodes, sb.nlog, sb.logstart, sb.inodestart, sb.bmapstart, IPB, BPB);

    return 0;
}
```

3. Write another program, say *printfs.c* and make other appropriate changes as you did in previous two labs to call **pfs()**.

- a. Implementing *printfs.c*

```
#include "types.h"
#include "stat.h"
#include "user.h"
#include "fcntl.h"
#include "fs.h"

int main(int argc, char *argv[])
{
    pfs();
    exit();
}
```

- b. Adding *printfs.c* to the Makefile

```
UPROGS=\
    _cat\
    _echo\
    _forktest\
    _grep\
    _init\
    _kill\
    _ln\
    _ls\
    _mkdir\
    _rm\
    _sh\
    _stressfs\
    _usertests\
    _wc\
    _nice\
    _foo\
    _cp\
    _ps\
    _printfs\
    _zombie\

EXTRA=\
mkfs.c ulib.c user.h cat.c echo.c forktest.c grep.c kill.c\
ln.c ls.c mkdir.c rm.c stressfs.c usertests.c wc.c printf.c nice.c cp.c ps.c foo.c zombie.c\
printf.c umalloc.c\
README dot-bochsrc *.pl toc.* runoff runoff1 runoff.list\
.gdbinit.tmpl gdbutil\
```

### c. Output

```
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ printf
sb: size  nblocks  ninodes nlog logstart inodestart bmap-start Inodes-per-block Bitmap-bits-per-block
    1000      941      200    30      2         32        58         8        4096
$ █
```

4. Add to your program and function call that prints out the free blocks and allocated blocks of the system as well as the number of free blocks and the number of allocated blocks in the system. The output of your program could look like the following:

#### a. Modified *pfs.c* function in *fs.c*

```
int
pfs()
{
    int b, bi, m;
    struct buf *bp;
    int countt = 0, countf = 0, counta = 0, bn = 0, nb = 0;
    bp = 0;
    cprintf("\nFree blocks: \n");
    for(b = 0; b < sb.size; b += BPB){
        nb++;
        bp = bread(1, BBLOCK(b, sb));
        for(bi = 0; bi < BPB && b + bi < sb.size; bi++){
            m = 1 << (bi % 8);
            ++countt;
            if((bp->data[bi/8] & m) == 0) { //checks to see if block is free
                if((countf % 16) == 0) cprintf("\n");
                cprintf(" %d ", bn);
                ++countf;
            }
            bn++;
        }
        brelse(bp);
    }
    cprintf("\nTotal free blocks = %d", countf);
    cprintf("\nAllocated blocks:\n");
    bn = 0;

    for(b = 0; b < sb.size; b += BPB) {
        bp = bread(1, BBLOCK(b, sb));
        for (bi = 0; bi < BPB && b + bi < sb.size; bi++) {
            m = 1 << (bi % 8);
            if((bp->data[bi/8] & m) != 0) { //checks to see if block is allocated
                ++counta;
                if((counta % 16) == 0) {
                    cprintf("\n");
                }
                cprintf(" %d", bn);
            }
            bn++;
        }
        brelse(bp);
    }
    cprintf("\nTotal allocated blocks = %d\n", counta);
    cprintf("\nTotal blocks = %d\n", countt);
    return 0;
}
"fs.c" 718L, 16925C written
```

---

## b. Output

```
xv6...
cpu1: starting 1
cpu0: starting 0
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32
init: starting sh
$ printf
```

Free blocks:

```
712 713 714 715 716 717 718 719 720 721 722 723 724 725 726 727
728 729 730 731 732 733 734 735 736 737 738 739 740 741 742 743
744 745 746 747 748 749 750 751 752 753 754 755 756 757 758 759
760 761 762 763 764 765 766 767 768 769 770 771 772 773 774 775
776 777 778 779 780 781 782 783 784 785 786 787 788 789 790 791
792 793 794 795 796 797 798 799 800 801 802 803 804 805 806 807
808 809 810 811 812 813 814 815 816 817 818 819 820 821 822 823
824 825 826 827 828 829 830 831 832 833 834 835 836 837 838 839
840 841 842 843 844 845 846 847 848 849 850 851 852 853 854 855
856 857 858 859 860 861 862 863 864 865 866 867 868 869 870 871
872 873 874 875 876 877 878 879 880 881 882 883 884 885 886 887
888 889 890 891 892 893 894 895 896 897 898 899 900 901 902 903
904 905 906 907 908 909 910 911 912 913 914 915 916 917 918 919
920 921 922 923 924 925 926 927 928 929 930 931 932 933 934 935
936 937 938 939 940 941 942 943 944 945 946 947 948 949 950 951
952 953 954 955 956 957 958 959 960 961 962 963 964 965 966 967
968 969 970 971 972 973 974 975 976 977 978 979 980 981 982 983
984 985 986 987 988 989 990 991 992 993 994 995 996 997 998 999
Total free blocks = 288
```

Allocated blocks:

```
01234567891011121314
15161718192021222324252627282930
31323334353637383940414243444546
47484950515253545556575859606162
63646566676869707172737475767778
79808182838485868788899091929394
9596979899100101102103104105106107108109110
111112113114115116117118119120121122123124125126
127128129130131132133134135136137138139140141142
143144145146147148149150151152153154155156157158
159160161162163164165166167168169170171172173174
175176177178179180181182183184185186187188189190
191192193194195196197198199200201202203204205206
207208209210211212213214215216217218219220221222
223224225226227228229230231232233234235236237238
239240241242243244245246247248249250251252253254
255256257258259260261262263264265266267268269270
271272273274275276277278279280281282283284285286
287288289290291292293294295296297298299300301302
303304305306307308309310311312313314315316317318
319320321322323324325326327328329330331332333334
335336337338339340341342343344345346347348349350
351352353354355356357358359360361362363364365366
367368369370371372373374375376377378379380381382
383384385386387388389390391392393394395396397398
399400401402403404405406407408409410411412413414
415416417418419420421422423424425426427428429430
431432433434435436437438439440441442443444445446
447448449450451452453454455456457458459460461462
463464465466467468469470471472473474475476477478
479480481482483484485486487488489490491492493494
495496497498499500501502503504505506507508509510
511512513514515516517518519520521522523524525526
527528529530531532533534535536537538539540541542
543544545546547548549550551552553554555556557558
559560561562563564565566567568569570571572573574
575576577578579580581582583584585586587588589590
591592593594595596597598599600601602603604605606
607608609610611612613614615616617618619620621622
623624625626627628629630631632633634635636637638
639640641642643644645646647648649650651652653654
655656657658659660661662663664665666667668669670
671672673674675676677678679680681682683684685686
687688689690691692693694695696697698699700701702
703704705706707708709710711
Total allocated blocks = 712
```

Tall blocks = 1000